

Homework 3 (Python version)

ME570 - Prof. Tron

2021-10-25

In this homework we will first introduce a special type of environment, a *sphere world*. You will then implement a potential-based planner (with two shapes for the attractive potential), a CBF-based planner, and a potential-based Inverse Kinematics solver for the two-link manipulator.

General instructions

Programming For your convenience, together with this document, you will find a **zip** archive containing Python files with stubs for each of the questions in this assignment; each stub contains an automatically generated description and header of the function or class. You will have to complete these files with the requested code. The goal of this files is to save you a little bit of time, and to avoid misspellings in the function or argument names. The files for the parts marked as **provided** (see also the *Grading* paragraph below) contain already the body of the function.

Homework help For best coding practices, please refer to the guidelines on Blackboard under Class content/Programming Tips & Tricks/Python. For questions specific to the content of the homework, please post on the Blackboard discussion board.

Homework report Along the programming of the requested functions, prepare a PDF report containing one or two sentences of comments for each question marked as **report**, and including: embedded figures and outputs that are representative of those generated by your code. Include comments on the questions marked as **code** only to explain any difficulty you might have encountered.

A small amount of *beauty points* are dedicated to reward reports that present their content in a professional way (see the *Grading criteria* section in the syllabus).

Analytical derivations To include the analytical derivations in your report you can type them in L^AT_EX (preferred method), any equation editor or clearly write them on paper and use a scanner (least preferred method).

Submission

The submission will be on Gradescope through four separate assignments: two for the questions marked as **code**, and one for those marked as **report**, and one for providing feedback. Further details are explained below. You can submit as many times as you would like, up to the assignment deadline. Each question is worth 1 point unless otherwise noted. Please refer to the Syllabus on Blackboard for late homework policies.

Report Upload the PDF of your report, and then indicate, for each question marked as **report**, on which page it is answered (just follow the Gradescope interface). Note that some of the questions marked as **report** might include a coding component, which however will be evaluated from the output figures you include in the report. In general, these questions are intended as checkpoints for you to visually check the results of your functions.

Code questions Upload all the necessary `.py` files, both those written by you, and those provided with the assignment. You will see *two* assignments on Gradescope. In the one marked *Python files*, you will submit the `.py` files directly; Gradescope will run one test checking for completeness, and one test checking the style (using Pylint); these automated tests will not check the correctness of your answers. In the other assignment marked *Python PDF listing*, please submit a PDF containing a print-out of the contents of the `.py` files (see the footnote¹ for how to generate such file). I will use this to manually grade the code. However, you can use the output of the test functions to judge if your code gives correct results or not.

Optional and provided questions. Questions marked as **optional** are provided just to further your understanding of the subject, and not for credit (if submitted, I will provide comments but it will not count toward your grade).

Hints

Some hints are available for some questions, and can be found at the end of the assignment (you are encouraged to try to solve the questions without looking at the hints first). If you use these hints, please state so in your report (your grading will not change based on this information, but it is a useful feedback for me).

Use of external libraries and toolboxes You are **not allowed** to use functions or scripts from external libraries or toolboxes (e.g., mapping toolbox), unless specifically instructed to do so (e.g., CVX).

Problem 1: Drawing and collision checking for spheres

In this problem you will write functions that are similar to those in Problem 1 of Homework 1, but applied to 2-D spheres (i.e., circles).

Question provided 1.1. A class for spheres.

File name: `me570_geometry.py`

Class name: `Sphere`

Description: Class for plotting and computing distances to spheres (circles, in 2-D).

¹The provided `.zip` file includes a `pygments_homework3.py` file that you can use to generate a HTML file with all the necessary pretty-printed listings via the command `python pygments_homework3.py`. The script requires the `pygments` Python package. You can then generate a PDF file using a browser and its “Print to pdf” functionality.

Method name: `__init__`

Description: Save the parameters describing the sphere as internal attributes.

Input arguments


- `center` (dim. $[2 \times 1]$): the 2-D coordinate of the center of the sphere;
- `radius` : a scalar whose absolute value is equal to the geometric radius of the sphere, and the sign indicates whether the obstacle should be interpreted as filled-in (`radius` > 0) or hollow (`radius` < 0);
- `distance_influence` : a scalar containing the *influence* distance of the sphere (the exact meaning of this will become clear after we talk about potential-based planning).

Figure 1 demonstrates the meaning of these quantities.


Question provided 1.2. Implement a function that draws a sphere.

Class name: `Sphere`

File name: `me570_geometry.py`



(a) Filled sphere (`radius` > 0)



(b) Hollow sphere (`radius` < 0)

Figure 1: Examples of spheres.

Method name: `plot`

Description: This function draws the sphere (i.e., a circle) of the given radius, and the specified color, and then draws another circle in gray with radius equal to the distance of influence.

Input arguments

- `color` (dim. $[1 \times 1]$, type `string`): a color specification string (e.g., `'b'`, `'r'`, etc.)

Question code 1.1. Implement the following function.

Class name: `Sphere`

File name: `me570_geometry.py`

Method name: `distance`

Description: Computes the signed distance between points and the sphere, while taking into account whether the sphere is hollow or filled in.

Input arguments

- `points` (dim. $[2 \times \text{nb_points}]$): an array of 2-D points.

Output arguments

- `d_points_sphere` (dim. $[1 \times \text{nb_points}]$): the distance between each point and the surface of the sphere. The distance is negative if `points` is inside the obstacle (i.e., inside the sphere for filled-in obstacles, and outside for hollow obstacles), and positive otherwise.

Requirements: Remember that the radius of the sphere is negative for hollow spheres.

Question code 1.2. Implement a function to compute the gradient of the sphere.

Class name: `Sphere`

File name: `me570_geometry.py`

Method name: `distance_grad`

Description: Computes the gradient of the signed distance between points and the sphere, consistently with the definition of `Sphere.distance` (`_`).

Input arguments

- `sphere` (dim. $[1 \times 1]$, type `struct`): a structure, as described above, defining a sphere.
- `points` (dim. $[2 \times \text{nb_points}]$): an array of 2-D points.

Output arguments

- `grad_d_points_sphere` (dim. $[2 \times \text{nb_points}]$): the gradient of the distance. If a point corresponds to the center of the sphere, return the zero vector.

Question optional 1.1. Implement a function that shows collision checks with a sphere.

File name: `me570_hw3.py`

`sphere_testCollision` (`_`)

Description: Generates one figure with a sphere (with arbitrary parameters) and `nb_points=100` random points that are colored according to the sign of their distance from the sphere (red for negative, green for positive). Generates a second figure in the same way (and the same set of points) but flipping the sign of the radius `r` of the sphere. For each sampled point, plot also the result of the output `pointsSphere`.

Red and green points represent locations that are, respectively, in collision and not in collision with the spherical obstacle.

The world

The goal of this question is to prepare and visualize the *sphere world* workspace that will be used in the other problems. The workspace is described by the data stored in the file `sphereworld.mat` provided with this homework.

File name: `me570_potential.py`

Class name: `SphereWorld`

Description: Class for loading and plotting a 2-D sphereworld.

Method name: `__init__`

Description: Load the sphere world from the provided file `sphereworld.mat`, and sets the following attributes:

- `world`: a `nb_spheres` list of `Sphere` objects defining all the spherical obstacles in the sphere world.
- `x_start`, a $[2 \times \text{nb_start}]$ array of initial starting locations (one for each column).
- `x_goal`, a $[2 \times \text{nb_goal}]$ vector containing the coordinates of different goal locations (one for each column).

Question provided 1.3. Implement a function to visualize a sphere world

Class name: `SphereWorld`

File name: `me570_potential.py`

Method name: `plot`

Description: Uses `Sphere.plot(_)` to draw the spherical obstacles together with a `*` marker at the goal location.

Use this function to visualize the contents of the file.

Problem 2: The potential-based planner

In this problem, you will implement and test the path planning strategy based on attractive-repulsive potential. Different functions will use similar dictionaries, which are explained here for ease of reference:

- A dictionary `potential` with elements:
 - `potential['x_goal']`: a $[2 \times 1]$ vector with the x, y coordinates of the goal location.
 - `potential['repulsive_weight']`: a scalar containing the weight of the attractive term with respect to the repulsive term.
 - `potential['shape']`: a string that specifies the shape of attractive potential. It can be equal to `'conic'` or `'quadratic'`.

How these structures are actually used is explained in the questions below.

In the following, we use the shorthand notation $d_i(x)$ to denote the distance between the point x and the sphere (defined as seen in class).

The attractive potential we use for this assignment is given by:

$$U_{\text{attr}}(x) = d^p(x, x_{\text{goal}}) = \|x - x_{\text{goal}}\|^p, \quad (1)$$

where p is a parameter to distinguish between conic and quadratic potentials, and its corresponding gradient is:

$$\nabla U_{\text{attr}}(x) = p d^{p-1}(x, x_{\text{goal}}) \frac{x - x_{\text{goal}}}{\|x - x_{\text{goal}}\|} = p d^{p-2}(x, x_{\text{goal}})(x - x_{\text{goal}}), \quad (2)$$

The repulsive potential we use for this assignment is given by

$$U_{\text{rep},i}(x) = \begin{cases} \frac{1}{2} \left(\frac{1}{d_i(x)} - \frac{1}{d_{\text{influence}}} \right)^2 & \text{if } 0 < d_i(x) < d_{\text{influence}}, \\ 0 & \text{if } d_i(x) > d_{\text{influence}}, \\ \text{undefined} & \text{otherwise,} \end{cases} \quad (3)$$

and the corresponding gradient is

$$\nabla U_{\text{rep},i}(x) = \begin{cases} - \left(\frac{1}{d_i(x)} - \frac{1}{d_{\text{influence}}} \right) \frac{1}{d_i(x)^2} \nabla d_i(x) & \text{if } 0 < d_i(x) < d_{\text{influence}}, \\ 0 & \text{if } d_i(x) > d_{\text{influence}}, \\ \text{undefined} & \text{otherwise.} \end{cases} \quad (4)$$

Question provided 2.1.

Class name: `RepulsiveSphere`

File name: `me570_potential.py`

Description: Repulsive potential for a sphere

method name: `__init__`

Description: Save the arguments to internal attributes

Input arguments

- `sphere` (type `Sphere`): an object describing a sphere as described above.

Question code 2.1. This function evaluates the repulsive potential.

Class name: `RepulsiveSphere`

File name: `me570_potential.py`

Method name: `eval`

Description: Evaluate the repulsive potential from `sphere` at the location $x = \text{x_eval}$. The function returns the repulsive potential as given by (3).

Input arguments

- `x_eval` (dim. $[2 \times 1]$): the location x at which to evaluate the total repulsive potential.

Output arguments

- `u_rep` (dim. $[1 \times 1]$): the value of U_{rep} evaluated at $x = \text{xEval}$.

Use the value *Not-a-Number* (`math.nan` in Python) for the case where the potential is undefined.

Question code 2.2. For this question you need to implement the gradient for the repulsive potential for spheres.

Class name: RepulsiveSphere

File name: me570_potential.py

Method name: grad

Description: Compute the gradient of U_{rep} for a single sphere, as given by (4).

Input arguments

- `x_eval` (dim. $[2 \times 1]$): see input arguments for `RepulsiveSphere.eval(_)`.

Output arguments

- `grad_u_rep` (dim. $[2 \times 1]$): the gradient of U_{rep} evaluated at $x = \text{xEval}$.

Requirements: This function should use `self.sphere.distance_grad(_)`

Use the value *Not-a-Number* (`Nan` in Matlab) for the case where the potential is undefined.

Question provided 2.2.

Class name: Attractive

File name: me570_potential.py

Description: Repulsive potential for a sphere

method name: `__init__`

Description: Save the arguments to internal attributes

Input arguments

- `potential` (dim. $[1 \times 1]$, type `struct`): the dictionary with fields `'x_goal'`, `'shape'` and `'repulsive_weight'` previously described (although the value of the field `'repulsive_weight'` is not used in this class).

Question code 2.3. Implement the attractive potential

Class name: Attractive

File name: me570_potential.py

Method name: eval

Description: Evaluate the attractive potential ∇U_{attr} at a point `xEval` with respect to a goal location `potential.xGoal` given by the formula: If `potential.shape` is equal to `'conic'`, use $p = 1$. If `potential.shape` is equal to `'quadratic'`, use $p = 2$.

Input arguments

- `x_eval` (dim. $[2 \times 1]$): the location x at which to evaluate the potential.

Output arguments

- `u_attr` (dim. $[1 \times 1]$): the value of U_{attr} evaluated at $x = \text{xEval}$.

Question code 2.4 (2 points). Implement a function to compute the gradient of the attractive potential.

Class name: `Attractive`

File name: `me570_potential.py`

Method name: `grad`

Description: Evaluate the gradient of the attractive potential ∇U_{attr} at a point `xEval`. The gradient is given by the formula If `potential['shape']` is equal to `'conic'`, use $p = 1$; if it is equal to `'quadratic'`, use $p = 2$.

Input arguments

- `x_eval` (dim. $[2 \times 1]$): the location x at which to evaluate the potential.

Output arguments

- `grad_u_attr` (dim. $[2 \times 1]$): the value of U_{attr} evaluated at $x = \text{xEval}$.

Question provided 2.3. The following is a utility function.

File name: `me570_geometry.py`

Method name: `field_plot_threshold`

Description: The function evaluates the function `f_handle` on points placed on a regular grid.

Input arguments

- `f_handle` (type `function`): an handle to a function that takes a vector `x` of dimension $[2 \times 1]$ as a single argument, and returns a scalar or a vector. An example of this function could be the function `norm(.)`.
- `threshold` **optional**, default `10`: if a vector `fEval`, obtained by evaluating `fHandle`, has `norm(fEval)>threshold`, then it is replaced by `fEval/norm(fEval)*threshold`. If omitted, `threshold = 10`.
- `nb_grid` **optional**, default `61`: size of the grid (number of points to the side).

Requirements: The function is build upon the function `Grid.eval(.)` from Homework 2.

Note that you can look at the source code of this function for a way to handle optional arguments.

Note on the use with functions that take extra parameters You can use `Field.plotThreshold(.)` also with functions that require additional parameters (e.g., `Potential.attractive(.)`). This can be done by creating a *lambda function* where the additional parameters are fixed. The best way to understand this is with an example:

```
sphere=Sphere(2*np.ones((2,1)), -2,3)
f_handle= lambda point: sphere.distance(point)
field_plot_threshold(f_handle,10)
```

You need to type the example, copying from this PDF document will not work.

The first line of the example creates an object `sphere` from the class `Sphere`. The second line creates a lambda function with a single argument `point`, and stores it in the variable `f_handle`; since `sphere` does not appear in the argument list of the lambda function, it is “captured” from the environment (i.e., it is “remembered” inside the lambda function). Now `f_handle` can be used as a function with a single input argument (e.g., try `f_handle(np.zeros((2,1)))`). In the third line, this function is passed to `field_plot_threshold` for plotting. Note that this will still work even though `sphere` is not visible in scope inside of `field_plot_threshold` (as `sphere` is “remembered” inside the lambda function, this is not necessary).

Question provided 2.4. The following class combines the attractive and repulsive potentials into the total potential.

Class name: `Total`

File name: `me570_potential.py`

Description: Combines attractive and repulsive potentials

method name: `__init__`

Description: Save the arguments to internal attributes

Input arguments

- `world` (type `SphereWorld`): an instance of the class `SphereWorld` (see the corresponding description for details).
- `potential` (dim. $[1 \times 1]$, type `struct`): the dictionary with fields `'x_goal'`, `'shape'` and `'repulsive_weight'` previously described. This function uses the field `'repulsive_weight'`, and then the structure is passed to an instance of `Attractive`.

Question code 2.5. The following two functions should combine the attractive and repulsive potentials into the total potential.

Class name: `Total`

File name: `me570_potential.py`

Method name: `eval`

Description: Compute the function $U = U_{\text{attr}} + \alpha \sum_i U_{\text{rep},i}$, where α is given by the variable `potential.repulsive_weight`.

Input arguments

- `x_eval` (dim. $[2 \times 1]$): the point where to evaluate the total potential.

Output arguments

- `u_eval` : The value of the potential U evaluated at `x_eval`.

Class name: `Total`

File name: `me570_potential.py`

Method name: `grad`

Description: Compute the gradient of the total potential, $\nabla U = \nabla U_{\text{attr}} + \alpha \sum_i \nabla U_{\text{rep},i}$, where α is given by the variable `potential.repulsiveWeight`

Input arguments

- `x_eval` (dim. $[2 \times 1]$): the point where to evaluate the gradient of the total potential.

Output arguments

- `grad_u_eval` (dim. $[2 \times 1]$): The gradient of the potential, ∇U , evaluated at `x_eval`.

Question optional 2.1. Use the function `field_plotThreshold`(`_`) to visualize the attractive and repulsive potentials (first separately for a goal and obstacle, and then combined in the total potential), and their gradients, for different situations (try both conic and quadratic potentials, and both filled-in and hollow spherical obstacles). For the total potential, overlap the output of `field_plotThreshold`(`_`) on top of the output of `SphereWorld.plot`(`_`).

Question code 2.6. Implement a function for your main potential planner.

Class name: `Planner`

File name: `me570_potential.py`

Method name: `run`

Description: This function uses a given function (`planner_parameters['control']`) to implement a generic potential-based planner with step size `planner_parameters['epsilon']`, and evaluates the cost along the returned path. The planner must stop when either the number of steps given by `planner_parameters['nb_steps']` is reached, or when the norm of the vector given by `planner_parameters['control']` is less than $5 \cdot 10^{-3}$ (equivalently, `5e-3`).

Input arguments

- `x_start` (dim. $[2 \times 1]$): starting location.
- ~~`world` (type `SphereWorld`), `potential` (type `dict`): descriptions of the world and of the potential as previously described. These are passed to the function `planner_parameters['control']` :~~
- `planner_parameters` (type `struct`): a dict with four keys:
 - `'U'`: handle to the function for computing the value of the potential function;
 - `'control'`: handle to the function for computing the direction in which the planner should move (for gradient-based methods, this is the negative gradient of the potential function);
 - `'epsilon'`: value for the step size;
 - `'nb_steps'`: total number of steps.

Output arguments

- `x_path` (dim. $[2 \times \text{nb_steps}]$): sequence of locations generated by the planner, such that `x[:,1]=x_start` and `x[:,k+1]=x[:,k]+epsilon*controlCurrent`, where `controlCurrent` is obtained by evaluating `control` using `x[:,k]`, `world` and `potential` as arguments. If the planner stops for some `k<nb_steps`, the remaining entries should be filled with `NaN` (the special value *Not-a-Number*, which can be generated with the `NaN(.)` function in Matlab).
- `u_path` (dim. $[1 \times \text{nb_steps}]$): The array of values of the potential U evaluated at each of the points on the path (that is, `u_path[k]` is the value of U evaluated at `x_path[k]`). Again, use `NaN`'s to fill the array if the planner stops before `nb_steps`.

Question code 2.7. Implement a test function for the planner.

Class name: `Planner`

File name: `me570_potential.py`

Method name: `run_plot`

Description: This function performs the following steps:

- 1) Test
- 2) Loads the problem data via an object `world` of class `SphereWorld`.
- 3) For each goal location in `world.xGoal`:
 - (a) Uses the function `Sphereworld.plot` (.) to plot the world in a first figure.
 - (b) Create (C)reates the dictionary `planner_parameters` with the field `'U'` set to `Total.eval`, and the field `['control']` set to the negative of `Total.grad`.
 - (c) Calls the function `Planner.run` (.) with the problem data and the input arguments. The function needs to be called five times, using each one of the initial locations given in `x_start` (also provided in `sphereworld.mat`).
 - (d) After each call, plot the resulting trajectory superimposed to the world in the first subplot; in a second subplot, show `u_path` (using the same color and using the `semilogy` command).

Requirements: To avoid too much clutter, use separate pairs of figures for different goal locations (but keep all the different paths from the different starting points to the same goal location in the same pair of figures).

Note: Questions report 3.2–report 2.5 should be considered together.

Question report 2.1. Show the results of `Planner.run_plot` (.) for different interesting combinations of `potential['repulsive_weight']`, `potential['shape']`, and `plannerParameters.epsilon`. For the `plannerParameters.control` argument, pass a function computing the negative of the gradient, and for `plannerParameters.nb_steps`, use `100`. Start with α in the interval `0.01 – 0.1` and ϵ in the range `1e-3 – 1e-2`, and explore from there. Typically, adjustments in `repulsiveWeight` require subsequent adjustments in `epsilon`. For *every* case where the planner converges, add a plot where you zoom in closely around the final equilibrium. In your report, try to have all figures on the same one or two pages for ease of comparison. Hints are available for this question.

Question report 2.2. Use the function `field_plotThreshold` (.) to visualize the total potential U , and its gradient ∇U in two separate figures for each one of the combinations `'repulsive_weight'` and `'shape'` included in the previous question (report 3.2) (for this question, it is sufficient to consider only one of the two goals). Include the images in your report.

Question report 2.3. Comment on the effects of varying each one of the parameters `'repulsive_weight'` and `'epsilon'`. Explain why the planner behaves in the way you see, explicitly explaining what happens for the four cases where the two parameters are,

respectively, small/small, small/large, large/small, and large/large. Additionally, comment on whether the results you see in practice are consistent with what discussed in class.

Question report 2.4 (0.5 points). Comment on the relation between the value of the potential U toward the end of the iterations, versus the fact that the planner correctly succeeded or failed. Explain different reasons why the planner might fail, and why changing the various parameters can improve the situation.

Question report 2.5. What is the difference between the two goals included in the provided dataset? In relation to this, what is the effect of the parameter `'shape'`?

Problem 3: CLF-CBF formulation

This problem is similar to Problem 2, except that you will use a CLF-CBF formulation instead of a traditional gradient-based formulation. We will use U_{attr} in (1) as the CLF, and $d_i(x)$ (implemented in `Sphere.distance(.)`) as the CBF for each obstacle. Note that, since we control the position of our (point) robot directly, the dynamics of our system is input-affine with the form $\dot{x} = f(x) + g(x)u = u$, i.e., $f(x) = 0$ and $g(x) = I$, the identity matrix.

Preparation. First, we need to write the Quadratic Program (QP) that uses the minimum-effort CLF-CBF formulation (the supervisor version) to return a control $u^*(x)$ for any given x (see also the class notes). In particular, the QP will have the form

$$\begin{aligned} u^*(x) = \underset{u}{\operatorname{argmin}} \|u - \clubsuit\|^2 & \quad (5a) \\ \text{subject to } \spadesuit \leq 0, & \quad (\text{CBF constraint, } i\text{-th obstacle}) \quad (5b) \end{aligned}$$

Question report 3.1. In the report, write the expressions for \clubsuit (involving U_{attr}) and \spadesuit (involving $d_i(x)$). For this and the questions below, denote the constant appearing in the constraint (5b) as c_h . Please pay attention to the directions of the inequalities.

Question provided 3.1. As detailed in the Week 6 in-class activity, the provided function `qp_supervisor(.)` will solve the following QP.

$$\begin{aligned} u_{\text{opt}} = \underset{u \in \mathbb{R}^2}{\operatorname{argmin}} \|u - u_{\text{ref}}\|^2 & \quad (6) \\ \text{subject to } A_{\text{barrier}}u + b_{\text{barrier}} \leq 0 \end{aligned}$$

Question code 3.1. Write a function to compute u^* .

File name: `me570_potential.py`

Method name: `clfcbf_control`

Description: Compute u^* according to (5).

Input arguments

- `x_eval` (dim. $[2 \times 1]$), `world` (dim. $[\text{nb_spheres} \times 1]$, type `struct`), `potential` (dim. $[1 \times 1]$, type `struct`): same as in the class `Total(.)`.

Output arguments

- `u_opt` (dim. $[2 \times 1]$): The solution to the QP (5).

Requirements: This function should use `qp_supervisor` () from Question provided 3.1, by building A_{barrier} , b_{barrier} according to ♠ in (5b), using $c_h = \text{potential}[\text{'repulsive_weight'}]$.

Question report 3.2. Show the results of `potential_planner_runPlot` () for one combination of `['repulsive_weight']` and `['epsilon']` that makes the planner work reliably. For the argument `planner_parameters['control']` , use `@clfcfbf_control` , for `planner_parameters['nb_steps']` , use 20 (but increase to 25 if necessary), and for `potential['shape']` , use `'quadratic'` . In your report, try to have all figures on the same one or two pages for ease of comparison.

Question report 3.3. Use the function `field_plot_threshold` () to visualize the control field $u^*(x)$ for each one of the combinations `['repulsiveWeight']` included in the previous question (report 3.2). Make sure to superimpose the field on top of the `SphereWorld` map. Note that the computation of $u^*(x)$ will take significantly longer than just the gradient, hence you might want to reduce the size of the grid passed to `field_plotThreshold` () (e.g., use a 10×10 grid). Include the images in your report.

Question report 3.4. Comment on the trade-off between traditional gradient-based methods and a CLF-CBF formulation.

Question optional 3.1. Repeat the previous questions with `['shape']` set to `'conic'` .

Problem 4: Jacobian-based Inverse Kinematics (IK) for the two-link manipulator

In this problem, you will combine the attractive potential from Problem 2 with the Jacobian of the two-link manipulator from the previous homework to create a way to move the end effector of the manipulator to a specific configuration by acting on its joint angles. This procedure is equivalent to computing an inverse of the kinematic map, and it is therefore commonly called Inverse Kinematics.

Question report 4.1. Recall from Homework 2 that we denote the position of the end effector in the world reference frame as ${}^{\mathcal{W}}p_{\text{eff}}$. Based on the results of Homework 2, write an expression for the Jacobian matrix J such that $\frac{d}{dt}({}^{\mathcal{W}}p_{\text{eff}}) = J(\theta)\dot{\theta}$ (this was optional in Homework 2, but it is now mandatory).

Question code 4.1. Implement a function for computing the matrix J .

Class name: `TwoLink`

File name: `me570_robot.py`

Description: See description from previous homework assignments.

Method name: `jacobian_matrix`

Description: Compute the matrix representation of the Jacobian of the position of the end effector with respect to the joint angles as derived in Question report 4.1.

Input arguments

- `theta` (dim. $[2 \times 1]$) An array containing $\begin{bmatrix} \theta_1 \\ \theta_2 \end{bmatrix}$, the two joint angles for the two-link manipulator.

Output arguments

- `jtheta` (dim. $[2 \times 2]$): The matrix $J(\theta)$ defined in Question report 4.1.

If you already solved this question in Homework 2, please copy the relevant portion of your previous report here.

Question optional 4.1. Compare the results of `Jtheta*thetaDot`, where `[Jtheta]=twolink_jacobianMatrix(theta)`, with the results of `Twolink.jacobian(theta,thetaDot)`, for arbitrary values of `theta` and `thetaDot` (they should be the same).

Question provided 4.1. In this question you will adapt the potential planner from Problem 2 to work with the two-link manipulator, by pulling back the total potential and its gradient from \mathbb{R}^2 to the configuration space of the two-link manipulator.

Class name: `TwoLinkPotential`

File name: `me570_robot.py`

Description: Combines attractive and repulsive potentials

method name: `__init__`

Description: Save the arguments to internal attributes

Input arguments

- `world` (type `SphereWorld`): an instance of the class `SphereWorld` (see the corresponding description for details).
- `potential` (dim. $[1 \times 1]$, type `struct`): the dictionary with fields `'x_goal'`, `'shape'` and `'repulsive_weight'` previously described. This function uses the field `'repulsive_weight'`, and then the structure is passed to an instance of `Attractive`.

Question code 4.2. Now it is time to implement the evaluation of the total potential and of its gradient for the `TwoLinkPotential` class.

Class name: `TwoLinkPotential`

File name: `me570_robot.py`

Method name: `eval`

Description: Compute the potential U pulled back through the kinematic map of the two-link manipulator, i.e., $U({}^{\mathcal{W}}p_{\text{eff}}(\vec{\theta}))$, where U is defined as in Question code 2.5, and ${}^{\mathcal{W}}p_{\text{eff}}(\theta)$ is the position of the end effector in the world frame as a function of the joint angles $\vec{\theta} = \begin{bmatrix} \theta_1 \\ \theta_2 \end{bmatrix}$.

Input arguments

- `theta_eval` (dim. $[2 \times 1]$): the joint angles θ at which to evaluate the potential.

Output arguments

- `u_eval_theta` : The value of the potential U evaluated at `x_eval`.

Class name: `TwoLinkPotential`

File name: `me570_robot.py`

Method name: `grad`

Description: Compute the gradient of the potential U pulled back through the kinematic map of the two-link manipulator, i.e., $\nabla_{\vec{\theta}} U({}^{\mathcal{W}}p_{\text{eff}}(\vec{\theta}))$.

Input arguments

- `theta_eval` (dim. $[2 \times 1]$): an array containing the two joint angles $\begin{bmatrix} \theta_1 \\ \theta_2 \end{bmatrix}$ at which the gradient should be evaluated.

Output arguments

- `grad_u_eval_theta` (dim. $[2 \times 1]$): The gradient of the pulled-back potential, evaluated at `theta_eval`.

Note that for computing the gradient of the pulled-back function, you will need to use the Jacobian matrix given by `TwoLink.jacobianMatrix(_)` (see also the material from class).

Question code 4.3. Make a function `TwoLink.plot_animate(theta_path)` that uses the function `TwoLink.plot(_)` from Homework 2 to show the sequence of configurations on the path specified by `theta_path`. Optionally, use `Time.sleep(1./30)` to put a pause between each plot. In any case, ensure to make all the drawing overlap (so that you can see all the configurations in a single picture).

Question code 4.4. Implement a function that tests the potential planner applied to the two-link manipulator. This is similar to Question code 2.7, except that we are planning the trajectory of the end effector of the manipulator, instead of a free point.

Class name: `TwoLinkPotential`

File name: `me570_robot.py`

Method name: `run_plot`

Description: This function performs the same steps as `Planner.run_plot()` in Question code 2.7, except for the following:

- In step 3)c: `planner_parameters['U']` should be set to `@twolink_total`, and `planner_parameters['control']` to the negative of `@twolink_totalGrad`.
- In step 3)c: Use the contents of the variable `thetaStart` instead of `xStart` to initialize the planner, and use only the second goal `x_goal[:,1]`.
- In step 3)d: Use `Twolink.plotAnimate()` to plot a decimated version of the results of the planner. Note that the output `xPath` from `Potential.planner()` will really contain a sequence of joint angles, rather than a sequence of 2-D points. Plot only every 5th or 10th column of `xPath` (e.g., use `xPath(:,1:5:end)`). To avoid clutter, plot a different figure for each start.

Input arguments

- `plannerParameters` (type `struct`): same as in `Potential.planner()`, except that `planner_parameters['U']` and `planner_parameters['control']` can be left unset.

Question report 4.2. Show the results of `twolink_planner_runPlot()` for (one) combination of `['repulsive_weight']` and `epsilon` that makes the planner work reliably. For the argument `planner_parameters['nb_steps']`, use 400, and for `potential['shape']`, use 'quadratic'. Note that `planner_parameters['U']`, and `planner_parameters['control']` are set by the function, so they do not need to be set in the argument. In your report, try to have all figures on the same one or two pages for ease of comparison.

If it is too difficult to get the planner work reliably for all goal configurations, it is fine to include the results using different parameters for different goals, or just for a few out of the five goals.

Note that, in this problem, we are considering only collisions between the spheres and the end effector of the manipulator; we are not considering collisions between the spheres and the links of the manipulator; in other words, it is normal to have overlap between the manipulator and the spheres, as long as the end effector is not inside an obstacle.

Hint for question report 3.2: For the argument `planner_parameters['control']`, you can use `@(x) -potential_totalGrad(x,world,potential)`. If everything works, the plot of `URep` should be monotonically decreasing, at least for some very small `epsilon`, even if the planner does not get to the goal. As a rule of thumb, the values of `['repulsiveWeight']` and `epsilon` can be significantly different between the `'conic'` and `'quadratic'` shapes, and when `['repulsiveWeight']` is increased, the value of `epsilon` that works tends to decrease (given our discussion in class, can you explain this?) Generally, if the planner “goes crazy”, it is because either the repulsive gradient or the step size are too large. If the planner never enters the region of influence of the obstacles, the weight for the repulsive term is too large. If the planner does not make enough progress, the step size is too small.

Hint for question report 4.1: Note that expressions such as $av + bw$, where $a, b \in \mathbb{R}$ are scalars, and $v, w \in \mathbb{R}^2$ are vectors, can be equivalently written as the matrix-vector multiplication $\begin{bmatrix} v & w \end{bmatrix} \begin{bmatrix} a \\ b \end{bmatrix}$.